

Modern Cryptography

or: The Unofficial Notes on the Georgia Institute
of Technology's **CS6260**: *Applied Cryptography*



George Kudrayvtsev

george.k@gatech.edu

Last Updated: February 22, 2020
(through Module 6)

Many lattes were consumed throughout the making of these notes. ☺ If you found it useful and are in a generous mood, feel free to buy me another! Shoot me a donation on Venmo [@george_k_btw](#) or PayPal [kudrayvtsev@sbcglobal.net](#) with whatever this guide was worth to you.

If I've shared a class with you, I might've made a guide for it as well; check out my [other notes](#)!

Happy studying! ☺

0	Preface	3
1	Introduction	5
1.1	Overview	6
1.1.1	Symmetric Cryptography	6
1.1.2	Asymmetric Cryptography	7
1.1.3	Cryptanalysis	7
I	Symmetric Cryptography	9
2	Perfect Security	10
2.1	Notation & Syntax	10
2.2	One-Time Pads	11
2.2.1	The Beauty of XOR	12
2.2.2	Proving Security	12
3	Block Ciphers	15
3.1	Modes of Operation	16
3.1.1	ECB—Electronic Code Book	16
3.1.2	CBC—Cipher-Block Chaining	16
3.1.3	CBCC—Cipher-Block Chaining with Counter	17
3.1.4	CTR—Randomized Counter Mode	17
3.1.5	CTRC—Stateful Counter Mode	19
3.2	Security Evaluation	19
3.2.1	IND-CPA: Indistinguishability Under Chosen-Plaintext Attacks	20
3.2.2	IND-CPA-cg: A Chosen Guess	22
3.2.3	What Makes Block Ciphers Secure?	24
3.2.4	Random Functions	25
3.2.5	IND-CCA: Indistinguishability Under Chosen-Ciphertext Attacks	31
3.3	Summary	34
4	Message Authentication Codes	35
4.1	Notation & Syntax	35
4.2	Security Evaluation	36
4.2.1	UF-CMA: Unforgeability Under Chosen-Message Attacks	37
4.3	Mode of Operation: CBC-MAC	38
5	Hash Functions	40
5.1	Collision Resistance	41
5.2	Building Hash Functions	42
5.3	One-Way Functions	45
5.4	Hash-Based MACs	46
	Index of Terms	48

PREFACE

I read that Teddy Roosevelt once said, “Do what you can with what you have where you are.” Of course, I doubt he was in the tub when he said that.

— Bill Watterson, *Calvin and Hobbes*

Before we begin to dive into all things cryptography, I’ll enumerate a few things I do in this notebook to elaborate on concepts:

- An item that is **highlighted like this** is a “term;” this is some vocabulary that will be used and repeated regularly in subsequent sections. I try to cross-reference these any time they come up again to link back to its first defined usage; most mentions are available in the Index.
- An item that is **highlighted like this** is a “mathematical property;” such properties are often used in subsequent sections and their understanding is assumed there.
- An item in a **maroon box**, like...

BOXES: A Rigorous Approach

... this often represents fun and interesting asides or examples that pertain to the material being discussed. They are largely optional, but should be interesting to read and have value, even if it’s not immediately rewarding.

- An item in a **blue box**, like...

QUICK MAFFS: Proving That the Box Exists

... this is a mathematical aside; I only write these if I need to dive deeper into a concept that’s mentioned in lecture. This could be proofs, examples, or just a more thorough explanation of something

that might've been “assumed knowledge” in the text.

- An item in a green box, like...

DEFINITION 0.1: Example

... this is an important cryptographic definition. It will often be accompanied by a highlighted **term** and dive into it with some mathematical rigor.

[Linky](#) I also sometimes include margin notes like the one here (which just links back here) that reference content sources so you can easily explore the concepts further.

INTRODUCTION

Cryptography is hard.

— Anonymous

The purpose of a cryptographic scheme falls into three very distinct categories. A common metaphor used to explain these concepts is a legal document.

- **confidentiality** ensures content *secrecy*—that it can't be read without knowledge of some secret. In our example, this would be like writing the document in a language nobody except you and your recipient understand.
- **authenticity** guarantees content *authorship*—that its author can be irrefutably proven. In our example, this is like signing the original document in pen (assuming, of course, your signature was impossible to forge).
- **integrity** guarantees content *immutability*—that it has not been changed. In our example, this could be that you get an emailed copy of the signed document to ensure that its language cannot be changed post-signing.

Note that even though all three of these properties can go hand-in-hand, they are not mutually constitutive. You can have any of them without the others: you can just get a copy of an unsigned document sent to you in plain English to ensure its integrity later down the line.

Analysing any proposed protocol, handshake, or other cryptographic exchange through the lens of each of these principles will be enlightening. Not every scheme is intended to guarantee all three of them, and different methods are often combined to achieve more than one of these properties. In fact, cases in which only one of the three properties are necessary occur all the time. It's important to not make a cryptographic scheme more complicated than it needs to be to achieve a given purpose: complexity breeds bugs.

TRIVIA: Cryptography's Common Cast of Characters

It's really useful to anthropomorphize our discussion of the mathematical intricacies in cryptography. For that, we use a cast of characters whose names give us immediate insight into what we should expect from them.

- **Alice** and **Bob** are the most common sender-recipient pairing. They are generally acting in good faith and aren't trying to break the cryptographic scheme in question. If a third member is necessary, **Carol** will enter the fray (for consistency of the allusion to Lewis Carroll's *Alice in Wonderland* ☺).
- **Eve** and **Mallory** are typically the two members trying to break the scheme. Eve is a *passive* attacker (short for eavesdropper) that merely observes messages between Alice and Bob, whereas malicious Mallory is an *active* attacker who can capture, modify, and inject her own messages into exchanges between other members.

You can check out the [Wikipedia article](#) on the topic for more historic trivia and the full cast of characters.


1.1 Overview

There are a number of different categories of cryptography, each of which applies in different settings.

crypto graphy
secret writing

1.1.1 Symmetric Cryptography


The notion of **symmetric keys** comes from the fact that both the sender and receiver of encrypted information share the same secret key, K . This secret is the only thing that separates a viable receiver from an attacker.



Symmetric key algorithms are often very efficient and supported by hardware, but their fatal flaw lies in **key distribution**. If two parties need to share a secret without anyone else knowing, how do they get it to each other without already having a secure channel?

1.1.2 Asymmetric Cryptography

Asymmetric cryptography is built to solve this problem. Here, secrets are only known to one party; instead, a key (Pk, Sk) is broken into two mathematically-linked components. There is a **public key** that is broadcasted to the world, and a **private key** that must be kept secret.



Asymmetric cryptography is often used to mutually establish a secret key (without revealing it!) for use with symmetric key algorithms, since those are faster. It's a little counterintuitive: two strangers can “meet” and “speak” publicly, yet walk away having established a mutual secret.

The big assumption (which we'll attack later) is that there's a trusted public repository that lists everyone's authentic public keys; without this assumption, we're back at the key distribution problem.

1.1.3 Cryptanalysis

An important part of cryptography is *evaluation*: how good is a given scheme? There are a number of approaches that can help us attack a “secure” protocol:

trial-and-error The most straightforward approach is to identify a possible attack vector and try it out. If it results in a successful attack, the protocol needs to be patched, and you can try finding another vector. If it doesn't seem possible to find an attack vector, that does not make the protocol secure. And yet, how do we find another vulnerability...?

provable security The best way to ensure that a protocol cannot be exploited is to prove it mathematically. It typically relies on proof by contradiction: if an attack is found, then some underlying mathematical assumptions must be false. A good example of this is the factoring problem that underlies most asymmetric cryptography: we've demonstrated time and again the difficulty in factoring large prime numbers. If an efficient algorithm were to arise, a lot of things would go very wrong, very quickly.

PART I

SYMMETRIC CRYPTOGRAPHY

Contents

2	Perfect Security	10
2.1	Notation & Syntax	10
2.2	One-Time Pads	11
3	Block Ciphers	15
3.1	Modes of Operation	16
3.2	Security Evaluation	19
3.3	Summary	34
4	Message Authentication Codes	35
4.1	Notation & Syntax	35
4.2	Security Evaluation	36
4.3	Mode of Operation: CBC-MAC	38
5	Hash Functions	40
5.1	Collision Resistance	41
5.2	Building Hash Functions	42
5.3	One-Way Functions	45
5.4	Hash-Based MACs	46
	Index of Terms	48

PERFECT SECURITY

How long do you want these messages to remain secret? I want them to remain secret for as long as men are capable of evil.

— Neal Stephenson, *Cryptonomicon*

As mentioned in the [Introduction](#), algorithms in symmetric cryptography rely on all members having a shared secret. Let's cover the basic notation we'll be using to describe our schemes and then dive into some.

2.1 Notation & Syntax

For consistency, we'll need common conventions when referring to cryptographic primitives in a scheme.

A well-described symmetric encryption scheme covers the following:

- a **message space**, denoted as the \mathcal{MsgSp} or \mathcal{M} for short, describes the set of things which can be encrypted. This is typically unrestricted and (especially in the context of the digital world) includes anything that can be encoded as bytes.
- a **key generation algorithm**, \mathcal{K} , or the key space spanned by that algorithm, \mathcal{KeySp} , describes the set of possible keys for the scheme and how they are created. The space typically restricts its members to a specific length.
- a **encryption algorithm** and its corresponding **decryption algorithm** (\mathcal{E}, \mathcal{D}) describe how a message $m \in \mathcal{M}$ is converted to and from ciphertext. We will use the notation $\mathcal{E}(K, M)$ and $\mathcal{E}_K(M)$ interchangeably to indicate encrypting M using the key K (and similarly for \mathcal{D}).

A well-formed scheme *must* allow all valid messages to be en/decrypted. Formally, this means:

$$\forall m \in \mathcal{MsgSp}, \forall k \in \mathcal{KeySp} : \mathcal{D}(k, \mathcal{E}(k, m)) = m$$

An encryption scheme defines the message space and three algorithms: $(\mathcal{MsgSp}, \mathcal{E}, \mathcal{D}, \mathcal{K})$. The key generation algorithm often just pulls a random n -bit string from the entire $\{0, 1\}^n$ bit space; to describe this action, we use notation $K \xleftarrow{\$} \mathcal{KeySp}$. The encryption algorithm is often randomized (taking random input in addition to (K, M)) and stateful. We'll see deeper examples of all of these shortly.

2.2 One-Time Pads

A **one-time pad** (or OTP) is a very basic and simple way to ensure absolutely perfect encryption, and it hinges on a fundamental binary operator that will be at the heart of many of our symmetric encryption schemes in the future: **exclusive-or**.

Messages are encrypted with an equally-long sequence of random bits using XOR. With our notation, one-time pads can be described as such:

- the key space is all n -bit strings: $\mathcal{KeySp} = \{0, 1\}^n$
- the message space is the same: $\mathcal{MsgSp} = \{0, 1\}^n$
- encryption and decryption are just XOR:

$$\begin{aligned}\mathcal{E}(K, M) &= M \oplus K \\ \mathcal{D}(K, C) &= C \oplus K\end{aligned}$$

Can we be a little more specific with this notion of “perfect encryption”? Intuitively, a secure encryption scheme should reveal nothing to adversaries who have access to the ciphertext. Formally, this notion is called being Shannon-secure (and is also referred to as **perfect security**): the probability of a ciphertext occurring should be equal for any two messages.

DEFINITION 2.1: Shannon-secure

An encryption scheme, $(\mathcal{K}, \mathcal{E}, \mathcal{D})$, is **Shannon-secure** if:

$$\begin{aligned}\forall m_1, m_2 \in \mathcal{MsgSp}, \forall C : \\ \Pr[\mathcal{E}(K, m_1) = C] &= \Pr[\mathcal{E}(K, m_2) = C]\end{aligned}\tag{2.1}$$

That is, the probability of a ciphertext C must be equally-likely for any two messages that are run through \mathcal{E} .

Note that this doesn't just mean that a ciphertext occurs with equal probability for a *particular* message, but rather than *any* message can map to *any* ciphertext

with equal probability. It's often necessary but *not* sufficient to show that a specific message maps to a ciphertext with equal probability under a given key; additionally, it's necessary to show that all ciphertexts can be produced by a particular message (perhaps by varying the key).

Shannon security can also be expressed as a conditional probability,¹ where all messages are equally-probable (i.e. independent of being) given a ciphertext:

$$\forall m \in \mathcal{MsgSp}, \forall C : \\ \Pr[M = m | C] = \Pr[M = m]$$

Are one-time pads Shannon-secure under these definitions? Yes, thanks to XOR.

2.2.1 The Beauty of XOR

XOR is the only primitive binary operator that outputs 1s and 0s with the same frequency, and that's what makes it the perfect operation for achieving unpredictable ciphertexts.

Given a single bit, what's the probability of the input bit?

x	y	$x \oplus y$
1	1	0
1	0	1
0	1	1
0	0	0

Table 2.1: The truth table for XOR.

Suppose you have some $c = 0$ (where $c \in \{0, 1\}^1$); what was the input bit m ? Well it could've been 1 and been XOR'd with 1 OR it could've been 0 and been XOR'd with 0... Knowing c gives us no new information about the input: our guess is still as good as random chance ($\frac{1}{2} = 50\%$).

Now suppose you know that $c = 1$; are your odds any better? In this case, m could've been 1 and been XOR'd with 0 OR it could've been 0 and XOR'd with 1... Again, we can't do better than random chance.

By the very definition of being Shannon-secure, if we (as the attacker) can't do better than random chance when given a ciphertext, the scheme is perfectly secure.

2.2.2 Proving Security

So we did a bit of a hand-wavy proof to show that XOR is Shannon-secure, but let's be a little more formal in proving that OTPs are perfect as well.

¹ See *Definition 2.3* in Katz & Lindell, pp. 29, parts of which are available on [Google Books](https://books.google.com/books?id=78v3CQ3Ct8YC).

Theorem 2.1. *One-time pads are a perfect-security encryption scheme.*

Proof. We start by fixing an arbitrary n -bit ciphertext: $C \in \{0, 1\}^n$. We also choose a fixed n -bit message, $m \in \text{MsgSp}$. Then, what's the probability that a randomly-generated key $k \in \text{KeySp}$ will encrypt that message to be that ciphertext? Namely, what is

$$\Pr[\mathcal{E}(K, m) = C]$$

for our **fixed** m and C ? In other words, how many keys can turn m into C ?

Well, by the definition of the OTP, we know that this can only be true for a single key: $K = m \oplus C$. Well, since every bit counts, and the probability of a single bit in the key being “right” is $1/2$:

$$\begin{aligned} \Pr[\mathcal{E}(K, m) = C] &= \Pr[K = m \oplus C] \\ &= \underbrace{\frac{1}{2} \cdot \frac{1}{2} \cdot \frac{1}{2} \cdots}_{n \text{ times}} \\ &= \frac{1}{2^n} \end{aligned}$$

Note that this is true $\forall m \in \text{MsgSp}$, which fulfills the requirement for perfect security! Every message is equally likely to result in a particular ciphertext. \square

The problem with OTPs is that keys can only be used once. If we're going to go through the trouble of securely distributing OTPs,² we could just exchange the messages themselves at that point in time...

Let's look at what happens when we use the same key across two messages. From the scheme itself, we know that $C_i = K \oplus M_i$. Well if we also have $C_j = K \oplus M_j$, then:

$$\begin{aligned} C_i \oplus C_j &= (K \oplus M_i) \oplus (K \oplus M_j) \\ &= (K \oplus K) \oplus (M_i \oplus M_j) && \text{XOR is associative} \\ &= M_i \oplus M_j && a \oplus a = 0 \end{aligned}$$

Though this may seem like insignificant information, it actually can [reveal](#) quite a bit about the inputs, and eventually the entire key if it's reused enough times.

An important corollary of perfect security is what's known as the **impossibility result** (also referred to as the **optimality of the one-time pad** when used in that context):

² One could envision a literal physical pad in which each page contained a unique bitstring; if two people shared a copy of these pads, they could communicate securely until the bits were exhausted (or someone else found the pad). Of course, if either of them lost track of where they were in the pad, everything would be gibberish from then-on...

Theorem 2.2. *If a scheme is **Shannon-secure**, then the key space cannot be smaller than the message space. That is,*

$$|\text{KeySp}| \geq |\text{MsgSp}|$$

Proof. We are given an encryption scheme \mathcal{E} that is supposedly perfectly-secure. So we start by fixing a ciphertext with a specific key, ($K_1 \in \text{KeySp}$ and plaintext message, $m_1 \in \text{MsgSp}$):

$$C = \mathcal{E}(K_1, m_1)$$

We know for a fact, then, that at least one key exists that can craft C ; thus if we pick a key $K \in \text{KeySp}$ *at random*, there's a non-zero probability that we'd get C again:

$$\Pr[\mathcal{E}(K, m_1) = C] > 0$$

Suppose then there is a message $m_2 \in \text{MsgSp}$ which we can *never* get from decrypting C :

$$\Pr[\mathcal{D}(K, C) = m_2] = 0 \quad \forall K \in \text{KeySp}$$

By the correctness requirement of a valid encryption scheme, if a message can never be decrypted from a ciphertext, neither should that ciphertext result from an encryption of the message:

$$\Pr[\mathcal{E}(K, M) = C] = 0 \quad \forall K \in \text{KeySp}$$

However, that violates Shannon-secrecy, in which the probability of a ciphertext resulting from the encryption of *any* two messages is equal; that's not the case here:

$$\Pr[\mathcal{E}(K, M_1) = C] \neq \Pr[\mathcal{E}(K, M_2) = C]$$

Thus, our assumption is wrong: m_2 cannot exist! Meaning there *must* be some $K_2 \in \text{KeySp}$ that decrypts C : $\mathcal{D}(K_2, C) = M_2$. Thus, it must be the case that there are as many keys as there are messages. \square

Ideally, we'd like to encrypt long messages using short keys, yet this theorem shows that we cannot be perfectly-secure if we do so. Does that indicate the end of this chapter? Thankfully not. If we operate under the assumption that our adversaries are computationally-bounded, it's okay to relax the security requirement and make breaking our encryption schemes very, *very* unlikely. Though we won't have *perfect* secrecy, we can still do extremely well.

We will create cryptographic schemes that are computationally-secure under **Kerckhoff's principle**, which effectively states that *everything* about a scheme should be publicly-available except for the secret key(s).

BLOCK CIPHERS

THESE are one of the fundamental building blocks of symmetric cryptography: a **block cipher** is a tool for encrypting short strings. Well-known examples include AES and DES.

DEFINITION 3.1: Block Cipher

Formally, a block cipher is a **function family** that maps from a k -bit key and an n -bit input string to an n -bit output string:

$$\mathcal{E} : \{0, 1\}^k \times \{0, 1\}^n \mapsto \{0, 1\}^n$$

Additionally, $\forall K \in \{0, 1\}^k$, $\mathcal{E}_K(\cdot)$ is a **permutation** on $\{0, 1\}^n$. This means its inverse is well-defined; we denote it either as $\mathcal{E}_K^{-1}(\cdot)$ or the much more intuitive $\mathcal{D}_K(\cdot)$.

$$\begin{aligned} \forall M, C \in \{0, 1\}^n : \quad & \mathcal{E}_K(\mathcal{D}_K(C)) = C \\ & \mathcal{D}_K(\mathcal{E}_K(M)) = M \end{aligned}$$

In a similar vein, ciphertexts are unique, so $\forall C \in \{0, 1\}^n$, there exists a *single* M such that $C = \mathcal{E}_K(M)$.

MATH REVIEW: Functions

A function is **one-to-one** if every input value maps to a unique output value. In other words, it's when no two inputs map to the same output.

A function is **onto** if all of the elements in the range have a corresponding input. That is, $\forall y \exists x$ such that $f(x) = y$.

A function is **bijective** if it is both one-to-one and onto; it's a **permutation** if it maps a set onto itself. In our case, the set in question will typically be the set of all n -length bitstrings: $\{0, 1\}^n$.

3.1 Modes of Operation

Block ciphers are limited to encrypting an n -bit string, but we want to be able to encrypt arbitrary-length strings. A **mode of operation** is a way to combine block ciphers to achieve this goal. For simplicity, we'll assume that our arbitrarily-long messages are actually a multiple of a block length; if they weren't, we could just pad them, but we'll omit that detail for brevity.

3.1.1 ECB—Electronic Code Book

The simplest mode of operation is ECB mode, visually described in [Figure 3.1](#). Given an n -bit block cipher \mathcal{E} and a message of length nb , we could just encrypt it block by block. The decryption is just as easy, applying the inverse block cipher on each piece individually:

$$\begin{aligned}C[i] &= \mathcal{E}_K(M[i]) \\ M[i] &= \mathcal{D}_K(C[i])\end{aligned}$$




Figure 3.1: The ECB ciphering mode.

This mode of operation has a fatal flaw that greatly compromises its security: if two message blocks are identical, the ciphertexts will be as well. Furthermore, encrypting the same long message will result in the same long ciphertext. This mode of operation is never used, but it's useful to present here to highlight how we'll fix these flaws in later modes.

3.1.2 CBC—Cipher-Block Chaining

This mode of operation fixes both flaws in ECB mode and is usable in real symmetric encryption schemes. It introduces a random **initialization vector** or IV to keep each ciphertext random, and it chains the output of one block into the input of the next block.




Figure 3.2: The CBC ciphering mode.

Each message block is first chained via XOR with the previous ciphertext before being run through the encryption algorithm. Similarly, the ciphertext is run through the inverse then XOR'd with the previous ciphertext to decrypt. That is,

$$\begin{aligned} C[i] &= \mathcal{E}_K(M[i] \oplus C[i-1]) \\ M[i] &= \mathcal{D}_K(C[i]) \oplus C[i-1] \end{aligned}$$

(where the base case is $C[0] = IV$).

The IV can be sent out in the clear, unencrypted, because it doesn't contain any secret information in-and-of itself. If Eve intercepts it, she can't do anything useful with it; if Mallory modifies it, the decrypted plaintext will be gibberish and the recipient will know something is up.

However, if an initialization vector is **repeated**, there can be information leaked to keen attackers about the underlying plaintext.

3.1.3 CBCC—Cipher-Block Chaining with Counter

In this mode, instead of using a randomly-generated IV, a counter is incremented for each new *message* until it wraps around (which typically doesn't occur, consider 2^{128}). This counter is XOR'd with the plaintext to encrypt and decrypt:

$$\begin{aligned} C[i] &= \mathcal{E}_K(M[i] \oplus ctr) \\ M[i] &= \mathcal{D}_K(C[i]) \oplus ctr \end{aligned}$$

The downside of these two algorithms is not a property of security but rather of performance. Because every block depends on the outcome of the previous block, both encryption and decryption must be done in series. This is in contrast with...

3.1.4 CTR—Randomized Counter Mode

Like ECB mode, this encryption mode encrypts each block independently, meaning it can be parallelized. Unlike all of the modes we've seen so far, though, it does not




Figure 3.3: The CBC ciphering mode.

use a block cipher as its fundamental primitive.¹ Specifically, the encryption function does not need to be invertible. Whereas before we used \mathcal{E}_K as a mapping from a k -bit key and an n -bit string to an n -bit string (see Definition 2.2), we can now use a function that instead maps them to an m -bit string:

$$F : \{0, 1\}^k \times \{0, 1\}^l \mapsto \{0, 1\}^L$$




Figure 3.4: The CTR ciphering mode.

This is because both the encryption and decryption schemes use F_K directly. They rely on a randomly-generated value R as fuel, much like the IV in the CBC modes.² Notice that to decrypt $C[i]$ in Figure 3.4, one needs to first determine $F_K(R + i)$, then XOR that with the ciphertext to get $M[i]$. The plaintext is never run through the encryption algorithm at all; instead, $F_K(R + i)$ is used as a **one-time pad** for $M[i]$. That is,

$$\begin{aligned} C[i] &= M[i] \oplus \mathcal{E}_K(R + i) \\ M[i] &= C[i] \oplus \mathcal{E}_K(R + i) \end{aligned}$$

Note that in all of these schemes, the only secret is K (F and \mathcal{E} are likely standardized and known).

¹ In practice, though, F_K will generally be a block cipher. Even though this properly is noteworthy, it does not offer any additional security properties.

² In fact, I'm not sure why the lecture decides to use R instead of IV here to maintain consistency. They are mathematically the same: both R and IV are pulled from $\{0, 1\}^n$.

3.1.5 CTRC—Stateful Counter Mode

Just like CBC, this mode has a variant that uses a counter rather than a randomly-generated value.

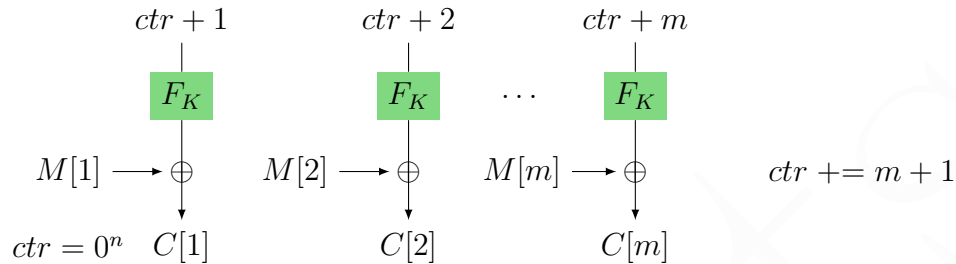


Figure 3.5: The CTRC ciphering mode.

3.2 Security Evaluation

Recall that we established that [Shannon-secure](#) schemes are impractical, and that we’re instead relying on adversaries being computationally bounded to achieve a reasonable level of security. To analyze our portfolio block ciphers, then, we need new definitions of this “computationally-bounded level of security.”

It’s easier to reverse the definition: a secure scheme is one that is not insecure. An insecure scheme allows a passive adversary that can see all ciphertexts do malicious things like learn the secret key or read any of the plaintexts. This isn’t rigorous enough, though: if the attacker can’t see any bits of the plaintext but can compute their sum, is that secure? What if they can tell when identical plaintexts are sent, despite not knowing their content?

There are plenty of possible information leaks to consider and it’s impossible to enumerate them all (especially when new attacks are still being discovered!). Dr. Boldyreva informally generalizes the aforementioned ideas:

Informally, an encryption scheme is secure if no adversary with “reasonable” resources who sees several ciphertexts can compute any³ partial information about the plaintexts, besides some a priori information.

Though this informality is not useful enough to prove things about encryption schemes we encounter, it’s enough to give us intuition on the formal definition ahead.

³ Any information *except* the length of the plaintexts; this knowledge is assumed to be public.


3.2.1 IND-CPA: Indistinguishability Under Chosen-Plaintext Attacks

It may be a mouthful, but the ability for a scheme to keep all information hidden when an attacker gets to feed their chosen inputs to it is key to a secure encryption scheme. **IND-CPA** is the formal definition of this attack.

We start with a fixed scheme and secret key, $\mathcal{SE} = (\text{KeySp}, \mathcal{E}, \mathcal{D})$; $K \xleftarrow{\$} \text{KeySp}$.

Consider an adversary \mathcal{A} that has access to an oracle. When they provide the oracle with a pair of equal-length messages, $m_0, m_1 \in \text{MsgSp}$, it outputs a ciphertext.

The oracle, called the “left-right encryption” oracle, chooses a bit $b \in \{0, 1\}^1$ to determine which of these messages to encrypt. It then passes m_b to the encryption function and outputs the ciphertext, $C = \mathcal{E}_K(m_b)$.



The adversary does not know the value of b , and thus does not know which of the messages was encrypted; it’s their goal to figure this out, given full access to the oracle. We say that an encryption scheme is secure if the adversary’s ability to determine which experiment the ciphertexts came from is no better than random chance.

DEFINITION 3.2: IND-CPA

A scheme \mathcal{SE} is considered secure under IND-CPA if an adversary’s **IND-CPA advantage**—the difference between their probability of guessing correctly and guessing incorrectly—is small (≈ 0):

$$\text{Adv}^{\text{ind-cpa}}(\mathcal{A}) = \Pr[\mathcal{A} \text{ guessed 0 for experiment 0}] - \Pr[\mathcal{A} \text{ guessed 0 for experiment 1}]$$

Since we are dealing with a “computationally-bounded” adversary, \mathcal{A} , we need to be cognizant about the real-world meaning behind resource usage. At the very least, we should consider the running time of our scheme and \mathcal{A} ’s attack. After all, if the encryption function itself takes an entire year, it’s likely unreasonable to give the attacker more than a few hundred tries at the oracle before they’re time-bound.

We should likewise be cognizant of how many queries the attacker makes and how long they are. We might be willing to make certain compromises of security if, for example, the attacker needs a 2^{512} -length message to gain an advantage.

With our new formal definition of security under our belt, let's take a crack at breaking the various [Modes of Operation](#) we defined. If we can provide an algorithm that demonstrates a reasonable advantage for an adversary that requires reasonable resources, we can show that a scheme is not secure under [IND-CPA](#).

Analysis of ECB

This was clearly the simplest and weakest of schemes that we outlined. The lack of randomness makes gaining an advantage trivial: the message can be determined by having a message with a repeating and one with a non-repeating plaintext.

ALGORITHM 3.1: A simple algorithm for breaking the ECB block cipher mode.

```

 $C_1 \parallel C_2 = \mathcal{E}_K(LR(0^{2n}, 0^n \parallel 1^n, b))$ 
if  $C_1 = C_2$  then
  | return 0
end
return 1

```

The attack can be generalized to give the adversary perfect knowledge for any input plaintext, and it leads to an important corollary.

Theorem 3.1. *Any deterministic, stateless encryption scheme cannot be IND-CPA secure.*

Proof. Under deterministic encryption, identical plaintexts result in identical ciphertexts. We can always craft an adversary with an advantage. First, we associate ciphertexts with plaintexts, then by the third message we can always determine which ciphertext corresponds to which input.

The proof holds for an arbitrary \mathcal{MsgSp} , we chose $\{0, 1\}^n$ in [algorithm 3.2](#) for convenience of representation. □

Analysis of CBC

Turns out, counters are far harder to “get right” relative to random [initialization vectors](#): their predictable nature means we can craft messages that are effectively deterministic by replicating the counter state. Namely, if we pre-XOR our plaintext with the counter, the first ciphertext block functions the same way as in ECB.

The first message lets us identify the counter value. The second message lets us craft

ALGORITHM 3.2: A generic algorithm for breaking deterministic encryption schemes.

```

 $C_1 = \mathcal{E}_K(LR(0^n, 0^n, b))$ 
 $C_2 = \mathcal{E}_K(LR(1^n, 1^n, b))$ 
// Given knowledge of these two, we can now always differentiate
// between them. We can repeat this for any  $m \in \text{MsgSp}$ .
 $C_3 = \mathcal{E}_K(LR(0^n, 1^n, b))$ 
if  $C_3 = C_1$  then
|   return 0
end
return 1

```

a “post-counter” message that will be equal to the third message.

ALGORITHM 3.3: A simple adversarial algorithm to break CBC mode.

```

// First, determine the counter (can't count on  $ctr = 0$ ).
 $C_0 \parallel C_1 = \mathcal{E}_K(LR(0^n, 1^n, b))$ 

// Craft a message that'll be all-zeros post-counter.
 $M_1 = 0^n \oplus (ctr + 1)$ 
 $C_2 \parallel C_3 = \mathcal{E}_K(LR(M_1, 1^n, b))$ 

// Craft it again, then compare equality.
 $M_3 = 0^n \oplus (ctr + 2)$ 
 $C_4 \parallel C_5 = \mathcal{E}_K(LR(M_3, 1^n, b))$ 
if  $C_3 = C_5$  then
|   return 0
end
return 1

```

3.2.2 IND-CPA-cg: A Chosen Guess

It turns out that the other modes of operation are provably IND-CPA secure *if* the underlying block cipher is secure. Before we dive into those proofs, though, let's define an alternative interpretation of the IND-CPA advantage; we will call this formulation **IND-CPA-cg**, for “chosen guess,” and it's shown visually in Figure 3.6. This formulation will be more convenient to use in some proofs.

In this version, the choice between left and right message is determined randomly at

the start and encoded within b . There is now only one experiment: if the attacker's guess matches ($b' = b$), the experiment returns 1.

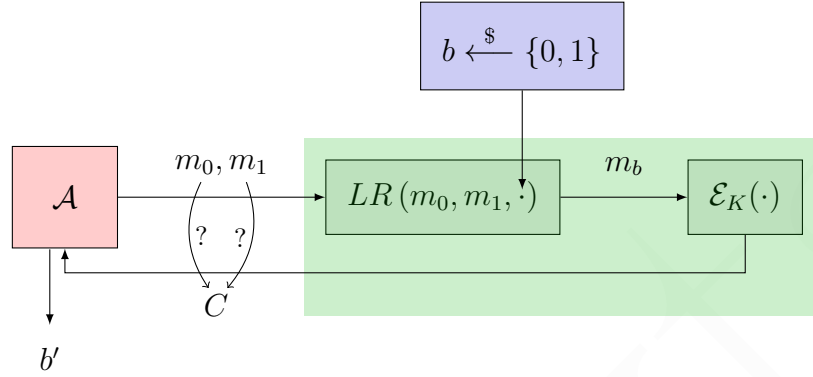


Figure 3.6: The “chosen guess” variant on IND-CPA security, where the attacker must guess a b' , and the experiment returns 1 if $b' = b$.

DEFINITION 3.3: IND-CPA-cg

A scheme \mathcal{SE} is still only considered secure under the “chosen guess” variant of IND-CPA if their **IND-CPA-cg advantage** is small; this advantage is now instead defined as:

$$\text{Adv}^{\text{ind-cpa-cg}}(\mathcal{A}) = 2 \cdot \Pr[\text{experiment returns 1}] - 1$$

The two variants on attacker advantage in Definition 2.3 and the new Definition 2.4 can be proven equal.

Claim 3.1. $\text{Adv}^{\text{ind-cpa}}(\mathcal{A}) = \text{Adv}^{\text{ind-cpa-cg}}(\mathcal{A})$ for some encryption scheme \mathcal{SE} .

Proof. The probability of the cg experiment being 1 (that is, the attacker guessing $b' = b$ correctly) can be expressed as conditional probabilities. Remember that $b \xleftarrow{\$} \{0, 1\}$ with uniformly-random probability.

$$\begin{aligned}
 \Pr[\text{experiment-cg returns 1}] &= \Pr[b = b'] \\
 &= \Pr[b = b' | b = 0] \Pr[b = 0] + \Pr[b = b' | b = 1] \Pr[b = 1] \\
 &= \Pr[b' = 0 | b = 0] \cdot \frac{1}{2} + \Pr[b' = 1 | b = 1] \cdot \frac{1}{2} \\
 &= \frac{1}{2} \cdot \Pr[b' = 0 | b = 0] + \frac{1}{2} (1 - \Pr[b' = 0 | b = 1]) \\
 &= \frac{1}{2} + \frac{1}{2} (\Pr[b' = 0 | b = 0] - \Pr[b' = 0 | b = 1])
 \end{aligned}$$

Notice the expression in parentheses: the difference between the probability of the attacker guessing 0 correctly (that is, when it really is 0) and incorrectly. This is exactly [Definition 2.3](#): advantage under the normal IND-CPA definition! Thus:

$$\begin{aligned}
 \Pr[\text{exp-cg returns } 1] &= \frac{1}{2} + \frac{1}{2} \underbrace{\Pr[b' = 0 | b = 0] - \Pr[b' = 0 | b = 1]}_{\text{IND-CPA advantage}} \\
 &= \frac{1}{2} + \frac{1}{2} \text{Adv}^{\text{ind-cpa}}(\mathcal{A}) \\
 2 \cdot \Pr[\text{exp-cg returns } 1] - 1 &= \text{Adv}^{\text{ind-cpa}}(\mathcal{A}) \\
 \text{Adv}^{\text{ind-cpa-cg}}(\mathcal{A}) &= \text{Adv}^{\text{ind-cpa}}(\mathcal{A})
 \end{aligned} \tag{3.1}$$

□

3.2.3 What Makes Block Ciphers Secure?

We can now look into the inner guts of each [mode of operation](#) and classify some block ciphers as being “secure” under IND-CPA. Refer to [Definition 2.2](#) to review the mathematical properties of a [block cipher](#). Briefly, it is a function family with a well-defined inverse that maps every message to a unique ciphertext for a specific key.

First off, it’s important to recall that we expect attackers to be computationally-bounded to a reasonable degree. This is because block ciphers—and all symmetric encryption schemes, for that matter—are susceptible to an [exhaustive key-search](#) attack, in which an attacker enumerates every possible $K \in \text{KeySp}$ until they find the one that encrypts some known message to a known ciphertext. If we say $k = |\text{KeySp}|$, this obviously takes $O(k)$ time and requires on average 2^{k-1} checks, which is why k must be large enough for this to be infeasible.

FUN FACT: Historical Key Sizes

Modern block ciphers like [AES](#) use *at least* 128-bit keys (though 192 and 256-bit options are available) which is considered secure from exhaustive search.

The now-outdated block cipher [DES](#) (invented in the 1970s) had a 56-bit key space, and it had a particular property that could speed up exhaustive search by a factor of two. This means exhaustive key-search on DES takes $\approx 2^{54}$ operations which took about 23 years on a 25MHz processor (fast at the time of DES’ inception). By 1999, the key could be found in only 22 hours.

The improved triple-DES or 3DES block cipher used 112-bit keys, but it too was abandoned in favor of AES for performance reasons: doing three DES computations proved to be too slow for efficient practical use.

Obviously a block cipher is not necessarily secure just because exhaustive key-search is not feasible. We now aim to define some measure of security for a block cipher. Why can't we just use IND-CPA? Well a block cipher is deterministic *by definition*, and we saw in [Theorem 3.1](#), a deterministic scheme cannot be IND-CPA secure. Thus our definition is too strong! We need something weaker for block ciphers that is still lets us avoid all possible information leaks: nothing about the key, nothing about the plaintexts (or some property of the plaintexts), etc. should be revealed.

We will say that a block cipher is secure if its output ciphertexts “look” random; more precisely, it'd be secure if an attacker can't differentiate its output from a random function. Well... that requires a foray into random functions.

3.2.4 Random Functions

Let's say that $\mathcal{F}(l, L)$ defines the set of ALL functions that map from l -bit strings to L -bit strings:

$$\forall f \in \mathcal{F}(l, L) \quad f : \{0, 1\}^l \mapsto \{0, 1\}^L$$

A random function g is then just a random function from that set: $g(\cdot) \xleftarrow{\$} \mathcal{F}(l, L)$. Now because picking a function at random is the same thing as picking a bitstring at random,⁴ we can define g in pseudocode as a deterministic way of picking bitstrings (see [algorithm 3.4](#)).

ALGORITHM 3.4: $g(x)$, a random function.

```

Define a global array  $T$ 
if  $T[x]$  is not defined then
    |  $T[x] \xleftarrow{\$} \{0, 1\}^L$ 
end
return  $T[x]$ 

```


A function family is a **pseudorandom function** family (a PRF) if the input-output behavior of a random instance of the family is computationally indistinguishable from a truly-random function. This input-output behavior is defined by [algorithm 3.4](#) and is hidden from the attacker.

PRF Security


⁴ Any function we pick will map values to L -bit strings. Concatenating all of these output bitstrings together will result in some nL -bit string, with a $L2^L$ bitstring being longest if the function maps to *every* bitstring. Each chunk of this concatenated string is random, so we can just pick some random $L2^L$ -length bitstring right off the bat to pick g .

The security of a block cipher depends on whether or not an attacker can differentiate between it and a random function. Like with IND-CPA, we have two experiments. In the first experiment, the attacker gets the output of the block cipher E with a fixed $K \in \text{KeySp}$; in the second, it's a random function g chosen from the PRF matching the domain and range of E .

Experiment 1 (“real”)



Experiment 0 (“random”)



The attacker outputs their guess, b , which should be 1 if they think they're being fed outputs from the real block cipher and 0 if they think it's random. Then, their “advantage” is how much more often the attacker can guess correctly.

DEFINITION 3.4: Block Cipher Security

A block cipher is considered **PRF secure** if an adversary's **PRF advantage** is small (near-zero), where the advantage is defined as the difference in probabilities of the attacker choosing

$$\text{Adv}^{\text{prf}}(\mathcal{A}) = \Pr[\mathcal{A} \text{ returns 1 for experiment 1}] - \Pr[\mathcal{A} \text{ returns 1 for experiment 0}]$$

For AES, the PRF advantage is very small and its *conjectured* (not proven) to be PRF secure. Specifically, for running time t and q queries,

$$\text{Adv}_{\text{AES}}^{\text{prf}}(\mathcal{A}) \leq \underbrace{\frac{ct}{T_{\text{AES}}} \cdot 2^{-128}}_{\text{exhaustive key-search}} + \underbrace{q^2 \cdot 2^{-128}}_{\text{birthday paradox}} \quad (3.2)$$

We will use this as an upper bound when calling a function F PRF secure.

The second term comes from an interesting attack that can be applied to *all* block ciphers known as the birthday paradox. Recall that block ciphers are permutations, so for distinct messages, you always get distinct ciphertexts. The attack is simple: if you feed the PRF security oracle q distinct messages and get q distinct ciphertexts, you output $b = 1$; otherwise, you output $b = 0$. The only way you get $< q$ distinct ciphertexts is from a g that isn't one-to-one. The probability of this happening is the probability of [algorithm 3.4](#) picking the same bitstring for two xs , so 2^{-L} .

FUN FACT: The Birthday Paradox

Suppose you're at a house party with 50 other people. What're the chances that two people at that party share the same birthday? Turns out, it's really, *really* high: 97%, in fact!

The **birthday paradox** is the counterintuitive idea despite the fact that YOU are unlikely to share a birthday with someone, the chance of ANY two people sharing a birthday is actually extremely high.

In the context of cryptography, this means that as the number of outputs generated by a random function g increases, the probability of SOME two inputs resolving to the same output increases much faster.

Proving Security: CTRC Recall the **CTRC—Stateful Counter Mode** mode of operation. Armed with the new definition of **block cipher security**, we can prove that this mode is secure. We start by assuming that the underlying cryptographic primitives are secure (in this case, this is the block cipher). Then, we can leverage the **contrapositive** to prove it. Starting with the implication:

If a scheme \mathcal{T} is y -secure,
then a scheme \mathcal{S} is x -secure.

(for some fill-in-the-blank x, y s like “IND-CPA” or “PRF”), we instead aim to prove the contrapositive:

If a scheme \mathcal{S} is NOT x -secure,
then a scheme \mathcal{T} is NOT y -secure.

To bring this into context, we will show that our mode of operation \mathcal{S} being insecure implies that the block cipher \mathcal{T} is *not* PRF-secure. More specifically, using our definitions of security, we're trying to show that: there existing an x -adversary \mathcal{A} that can break \mathcal{S} implies that there exists a y -adversary \mathcal{B} that can break \mathcal{T} :

- We assume \mathcal{A} exists, then construct \mathcal{B} using \mathcal{A} .
- Then, we show that \mathcal{B} 's y -advantage is not “too small” if \mathcal{A} 's x -advantage is not “too small” (≈ 0).

LOGIC REVIEW: Contrapositive

The **contrapositive** of an implication is its inverted negation. Namely, for two given statements p and q :

$$\begin{array}{l} \text{if } p \implies q \\ \text{then } \neg q \implies \neg p \end{array}$$

With that in mind, let's prove CTRC's security. To be verbose, the statements we're aiming to prove are:

$$\underbrace{\text{the underlying blockcipher is secure}}_P \implies \underbrace{\text{CTRC is a secure mode of operation}}_Q$$

However, since we're approaching this via the contrapositive, we'll instead prove

$$\underbrace{\text{CTRC is not a secure mode of operation}}_{\neg Q} \implies \text{only when } \underbrace{\text{the underlying blockcipher is not secure}}_{\neg P}$$

Theorem 3.2. *CTRC is a secure mode of operation if its underlying block cipher is secure.*

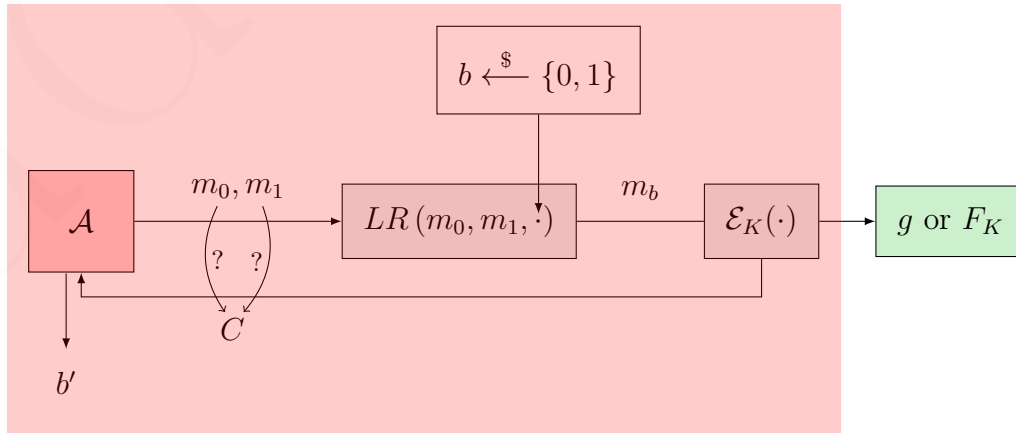
More formally, for any efficient adversary \mathcal{A} , $\exists \mathcal{B}$ with similar efficiency such that the IND-CPA advantage of \mathcal{A} under CTRC mode is less than double the PRF advantage of \mathcal{B} under a secure block cipher F :

$$\text{Adv}_{\text{CTRC}}^{\text{ind-cpa}}(\mathcal{A}) \leq 2 \cdot \text{Adv}_F^{\text{prf}}(\mathcal{B})$$

where we know an example of a secure block cipher $F = \text{AES}$ that any \mathcal{B} 's advantage will be very small (see (3.2)).

Proof. Let \mathcal{A} be an IND-CPA-cg adversary attacking CTRC. Then, we can present the PRF adversary \mathcal{B} .

- We construct \mathcal{B} so that it can act as the very left-right oracle that \mathcal{A} uses to query and attack the CTRC scheme.




adversary $\mathcal{B} \longrightarrow 1$ iff $b' = b$ else 0

- Namely, \mathcal{B} lets \mathcal{A} make oracle queries to CTRC until it guesses b correctly. This is valid because \mathcal{B} still delegates to a PRF oracle which is choosing between a random function g and the block cipher F_K (where K is still secret) for the actual block cipher; everything else is done exactly as described for CTRC in Figure 3.5.
- This construction lets us leverage the fact that \mathcal{A} knows how to break CTRC-encrypted messages, but we don't need to know how. For the pseudocode describing this process, refer to algorithm 3.5.

Now let's analyze \mathcal{B} , expressing its PRF advantage over F in terms of \mathcal{A} 's IND-CPA advantage over CTRC.⁵ The ability for \mathcal{B} to differentiate between F and some random function $g \in \text{Func}(\ell, L)$ depends *entirely* on \mathcal{A} 's ability to differentiate between CTRC with an actual block cipher F and a truly-random function g . Thus,

$$\begin{aligned}
 \text{Adv}_F^{\text{prf}}(\mathcal{B}) &= \Pr[\mathcal{B} \rightarrow 1 \text{ in } \text{Exp}_F^{\text{prf-0}}] - \Pr[\mathcal{B} \rightarrow 1 \text{ in } \text{Exp}_F^{\text{prf-1}}] && \text{definition} \\
 &= \Pr[\text{Exp}_{\text{CTRC}[F]}^{\text{ind-cpa-cg}} \rightarrow 1] - \Pr[\text{Exp}_{\text{CTRC}[g]}^{\text{ind-cpa-cg}} \rightarrow 1] && \mathcal{B} \text{ depends only on } \mathcal{A} \\
 &= \frac{1}{2} \cdot \text{Adv}_{\text{CTRC}[F]}^{\text{ind-cpa}}(\mathcal{A}) + \frac{1}{2} - \frac{1}{2} \cdot \text{Adv}_{\text{CTRC}[g]}^{\text{ind-cpa}}(\mathcal{A}) - \frac{1}{2} && \text{IND-CPA is equal to IND-CPA-cg via (3.1)}
 \end{aligned}$$

Next, we'll show that $\text{Adv}_{\text{CTRC}[g]}^{\text{ind-cpa}}(\mathcal{A}) = 0$. That is, we will show that \mathcal{A} has absolutely no advantage in breaking the scheme when using g —a truly-random function—as the block cipher. Consider the visualization of the CTRC scheme again:



Notice that the inputs to g are all distinct points, and by definition of a truly-random function its outputs are truly-random bitstrings. These are then XOR'd with messages... sound familiar? The outputs of g are distinct **one-time pads** and thus each $C[i]$ is **Shannon-secure**, meaning an advantage is simply impossible by definition.

⁵ The syntax $\mathcal{X} \rightarrow n$ means the adversary \mathcal{X} outputs the value n , and the syntax Exp_m^n refers to the experiment n under some parameter or scheme m , for shorthand.

The theorem claim can then be trivially massaged out:

$$\begin{aligned}
 \text{Adv}_F^{\text{prf}}(\mathcal{B}) &= \frac{1}{2} \cdot \text{Adv}_{\text{CTRC}[F]}^{\text{ind-cpa}}(\mathcal{A}) + \frac{1}{2} - \frac{1}{2} \cdot \underbrace{\text{Adv}_{\text{CTRC}[g]}^{\text{ind-cpa}}(\mathcal{A})}_{=0} - \frac{1}{2} \\
 &= \frac{1}{2} \cdot \text{Adv}_{\text{CTRC}}^{\text{ind-cpa}}(\mathcal{A}) \\
 2 \cdot \text{Adv}_F^{\text{prf}}(\mathcal{B}) &= \text{Adv}_{\text{CTRC}}^{\text{ind-cpa}}(\mathcal{A})
 \end{aligned}$$

□

ALGORITHM 3.5: Constructing an adversary \mathcal{B} that uses another adversary \mathcal{A} to break a higher-level symmetric encryption scheme.

Input: An adversary \mathcal{A} that executes oracle queries.

Result: 1 if \mathcal{A} succeeds in breaking the emulated scheme, 0 otherwise.

Let $g \xleftarrow{\$} F(\ell, L)$ where F is a **PRF**, which \mathcal{B} will use as a **block cipher**.

Let $\mathcal{E}_f(\cdot)$ be an encryption function that works like \mathcal{A} expects (for example, a CTRC scheme).

Choose a random bit: $b \xleftarrow{\$} \{0, 1\}^1$

repeat

 Get a query from \mathcal{A} , some (M_1, M_2)

$C \xleftarrow{\$} \mathcal{E}_f(M_b)$

 return C to \mathcal{A}

until \mathcal{A} outputs its guess, b'

return 1 iff $b = b'$, 0 otherwise

Proving Security: CTR Recall that the difference between CTRC (which we just proved was secure) and standard CTR is the use of a random IV rather than a counter (see [Figure 3.4](#)). It's also provably **PRF secure**, but we'll state its security level without proof:⁶

⁶ Feel free to refer to the [lecture video](#) to see the proof. In essence, the fact the value is chosen randomly means it's possible that for enough R s and M s there will be overlap for some $R_i + m$ and $R_j + n$. This will result in identical "one-time pads," though thankfully it occurs with a very small probability (it's related to the [birthday paradox](#)).

Theorem 3.3. *CTR is a secure mode of operation if its underlying block cipher is secure. More formally, for any efficient adversary \mathcal{A} , $\exists \mathcal{B}$ with similar efficiency such that:*

$$\text{Adv}_{CTR}^{\text{ind-cpa}}(\mathcal{A}) \leq 2 \cdot \text{Adv}_F^{\text{prf}}(\mathcal{B}) + \frac{\mu_A^2}{\ell 2^\ell}$$

where μ is the total number of bits \mathcal{A} sends to the oracle.

It's still secure because $\ell \geq 128$ for secure block ciphers, making the extra term near-zero. Proving bounds on security is very useful: we can see here that CTRC mode is better than CTR mode because there is no additional constant.

There is a similar theorem for CBC mode (see [Figure 3.2](#)), the last mode of operation whose security we haven't formalized.

Theorem 3.4. *CBC is a secure mode of operation if its underlying block cipher is secure. More formally, for any efficient adversary \mathcal{A} , $\exists \mathcal{B}$ with similar efficiency such that:*

$$\text{Adv}_{CBC}^{\text{ind-cpa}}(\mathcal{A}) \leq 2 \cdot \text{Adv}_F^{\text{prf}}(\mathcal{B}) + \frac{\mu_A^2}{n^2 2^n}$$

where μ is the total number of bits \mathcal{A} sends to the oracle.

We can see that $n^2 > \ell$ when comparing CBC to CTR, meaning the term will be smaller for the same μ . Thus, CTRC is more secure than CBC is more secure than CTR. The constant again comes from the [birthday paradox](#).

3.2.5 IND-CCA: Indistinguishability Under Chosen-Ciphertext Attacks

Is the intuition behind a scheme being both [IND-CPA](#) and [PRF](#) secure sufficient? Does IND-CPA take into account all of the possible attack vectors? Well, it limits attackers to choosing *plaintexts* and using only their ciphertext results to make learn information about the scheme and see ciphertexts. What if the attacker could instead attack a scheme by choosing *ciphertexts* and learn something about the scheme from the resulting plaintexts?

This isn't a far-fetched possibility,⁷ and it has historic precedent in being a viable attack vector. Since [IND-CPA](#) does not cover this vector, we need a stronger definition of security: the attacker needs more power. With [IND-CCA](#), the adversary \mathcal{A} has access to *two* oracles: the left-right encryption oracle, as before, and a decryption oracle.

⁷ Imagine reverse-engineering an encrypted messaging service like iMessage to fully understanding its encryption scheme, and then control the data that gets sent to Apple's servers to "skip" the encryption step and control the ciphertext directly. If you control both endpoints, you can see what the ciphertext decrypts to!

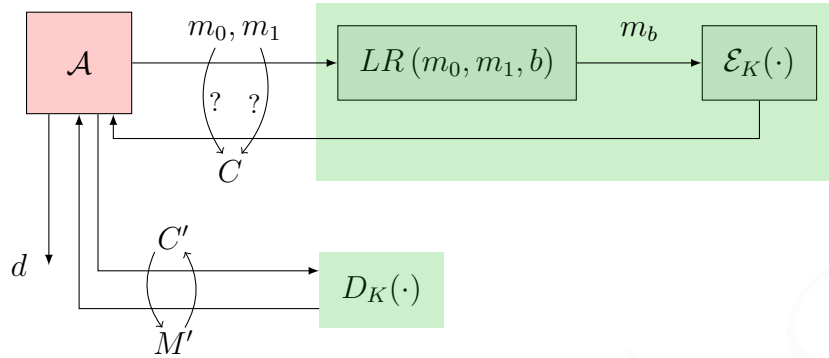


Figure 3.7: A visualization of the IND-CCA security definition. The adversary \mathcal{A} submits two messages, (m_0, m_1) , to an encryption oracle that (consistently) chooses one of them based on a bit b and returns m_b 's ciphertext. The adversary can also submit any C' that hasn't been submitted to LR to a decryption oracle and see the resulting plaintext.

The only restriction on the attacker is that they cannot query the decryption oracle on ciphertexts returned from the encryption oracle (obviously, that would make determining b trivial) (in Figure 3.7, this means $C \neq C'$). As before, a scheme is considered IND-CCA secure if an adversary's advantage is small.

DEFINITION 3.5: IND-CCA

A scheme \mathcal{SE} is considered secure under IND-CCA if an adversary's **IND-CCA advantage**—the difference between their probability of guessing correctly and guessing incorrectly—is small (≈ 0):

$$\text{Adv}^{\text{ind-cca}}(\mathcal{A}) = \Pr[\mathcal{A} \text{ guessed } 0 \text{ for experiment } 0] - \Pr[\mathcal{A} \text{ guessed } 0 \text{ for experiment } 1]$$

Note that since IND-CCA is stronger than IND-CPA, the former implies the latter. This is trivially-provable by reduction, so we won't show it here.

Unfortunately, none of our IND-CPA schemes are also secure under IND-CCA.

Analysis of CBC

Recall from Figure 3.2 the way that message construction works under CBC with random initialization vectors.

Suppose we start by encrypting two distinct, two-block messages. They don't have to be the ones chosen here, but it makes the example easier. We pass these to the left-right oracle:

$$IV \parallel c_1 \parallel c_2 \xleftarrow{\$} \mathcal{E}_K(LR(0^{2n}, 1^{2n}))$$

From these ciphertexts alone, we've already shown that the adversary can't determine which of the input messages was encrypted. However, suppose we send just the first chunk to the decryption oracle?

$$m = \mathcal{D}_K(IV \parallel c_1)$$

This is legal since it's not an *exact* match for any encryption oracle outputs. Well since our two blocks were identical, and c_2 has no bearing in the decryption of $IV \parallel c_1$ (again, refer to the visualization in [Figure 3.2](#)), the plaintext m will be all-zeros in the left case and all-ones in the right case!

It should be fairly clear that this is an efficient attack, and that the adversary's advantage is optimal (exactly 1). For posterity,

$$\begin{aligned} \text{Adv}_{\text{CBC}}^{\text{ind-cca}}(\mathcal{A}) &= \Pr[\mathcal{A} \rightarrow 0 \text{ for } \text{Exp}^{\text{ind-cca-0}}] - \Pr[\mathcal{A} \rightarrow 0 \text{ for } \text{Exp}^{\text{ind-cca-1}}] \\ &= 1 - 0 = \boxed{1} \end{aligned}$$

The attack time t is the time to compare n bits, it requires $q_e = q_d = 1$ query to each oracle, and message lengths of $\mu_e = 4n$ and $\mu_d = 2n$. Thus, CBC is not IND-CCA secure. \square

Almost identical proofs can be used to break both CTR and CTRC, our final bastions of hope in the [Modes of Operation](#) we've covered.

Analysis of CBC: Anotha' One

(or, *Kicking 'em While They're Down*)

We can break CBC (and the others) in a different way. This is included here to jog the imagination and offer an alternative way of thinking about showing insecurity under IND-CCA.

In this attack, one-block messages will be sufficient:

$$IV \parallel c_1 \xleftarrow{\$} \mathcal{E}_K(LR(0^n, 1^n))$$

This time, there's nothing to chop off. However, what if we try decrypting the ciphertext with a flipped IV?

$$m = \mathcal{D}_K(\overline{IV} \parallel c_1)$$

Well, according to [Figure 3.2](#), the output from the blockcipher will be XOR'd with the flipped IV, and thus result in a flipped message, so $m = \overline{0^n} = 1^n$ in the left case, and $m = 0^n$ in the right case!

Again, this is trivially computationally-reasonable (in fact, it's even *more* reasonable than before) and breaks IND-CCA security. \square

3.3 What Now?

We started off by enumerating a number of ways to create ciphertexts from plaintexts using [block ciphers](#). It was critical to follow that with several definitions of what “security” means, and we showed that some of the modes of operation (namely ECB and CBC) were not secure under [Definition 3.2](#), IND-CPA security. Then, we dug deeper to study the underlying block ciphers and what it meant for *those* to be PRF secure: it must be hard to differentiate them from a [random function](#) (see [Definition 3.3](#)). Finally, we gave the attacker more power under the last, strictest metric of security: IND-CCA, and showed that our remaining modes of operation (that is, CBC, CTR, and CTRC) broke under this adversarial scheme.


We definitely want a way to achieve IND-CCA security. Does hope remain? Thankfully, it does, but it will first require a foray into the other realms of cryptography: [integrity](#) and [authenticity](#).

MESSAGE AUTHENTICATION CODES

When in doubt, use brute force.

— Ken Thompson

DATA privacy and confidentiality is not the only goal of cryptography, and a good encryption method does not make any guarantees about anything beyond confidentiality. In the one-time pad (which is *perfectly* secure), an active attacker Mallory can modify the message in-flight to ensure that Alice receives something other than what Bob sent:



If Mallory knows that the first 8 bits of Bob’s message corresponds to the number of dollars that Alice needs to send Bob (and she does, according to Kerckhoff’s principle), such a manipulation will have catastrophic consequences for Alice’s bank account. Clearly, we need a way for Alice to know that a message came from Bob himself.

Let’s discuss ways to ensure that the recipient of a message can validate that the message came from the intended sender (authenticity) and was not modified on the way (integrity).

4.1 Notation & Syntax

A **message authentication code** (or MAC) is a fundamental primitive in achieving data authenticity under the symmetric cryptography framework. Much like in an encryption scheme, a well-defined MAC scheme covers the following:

- a **message space**, denoted as the \mathcal{MsgSp} or \mathcal{M} for short, describes the set of things which can be authenticated.

- a **key generation algorithm**, \mathcal{K} , or the key space spanned by that algorithm, KeySp , describes the set of possible keys for the scheme and how they are created.
- the MAC algorithm itself, \mathcal{MAC} (also called a **tagging** or **signing algorithm**) defines the way some $m \in \text{MsgSp}$ is authenticated and returns a tag.
- the MAC's corresponding **verification algorithm**, \mathcal{V}_F , describes how a message should be validated, given a (supposedly) authenticated message and its tag, outputting a Boolean value indicating their validity.

Succinctly, we say that $\Pi = (\mathcal{K}, \mathcal{MAC}, \mathcal{V}_F)$, and by definition

$$\forall k \in \text{KeySp}, \forall m \in \text{MsgSp} : \quad \mathcal{V}_F(k, m, \mathcal{MAC}(k, m)) = 1$$

If a MAC algorithm is deterministic, then \mathcal{V}_F does not need to be explicitly defined, since running the MAC on the message again and comparing the resulting tags is sufficient.

An important thing to remember in this chapter is that *we don't care about confidentiality*: the messages and their tags are sent in the clear. Our only concern is now **forging**—can Mallory pretend that a message came from Bob?

4.2 Security Evaluation

As before, with the [Security Evaluation](#) of a block cipher or its mode of operation, we need a way to model practical, strong adversaries and their attacks on MACs.

To start, we can imagine that an adversary can see some number of (message, tag) pairs. To mimic [IND-CCA](#), perhaps s/he can also force the tagging of messages and check the verification of specific pairs. Obviously, they shouldn't be able to compute the secret key, but more importantly, they should *never* be able to compose a message and tag pairing that is considered valid.

ATTACK VECTOR: Pay to (Re)Play

A **replay attack** is one where an adversary uses valid messages from the past that they captured to duplicate some action.


For example, imagine Bob sends an encrypted, authenticated message “You owe my friend Mallory \$5.” to Alice that everyone can see. Alice knows this message came from Bob, so she pays her dues. Then, Mallory decides to just... send Alice that message again! It's again deemed valid, and Alice

again pays her dues.

Protection against replay attacks requires some more-sophisticated construction of a secure scheme, so we'll ignore them for now as we discuss MAC schemes.

4.2.1 UF-CMA: Unforgeability Under Chosen-Message Attacks

Let's formalize these intuitions: the adversary \mathcal{A} is given access to two oracles that run the tagging and verification algorithms respectively, and s/he must output a message-tag pair (M, t) for which t is a valid tag for M (that is, $\mathcal{V}_F(K, M, t) = 1$) and M was never an input to the tagging oracle.¹



DEFINITION 4.1: UF-CMA

A message authentication code scheme Π is considered to be **UF-CMA** secure if the **UF-CMA advantage** of any adversary \mathcal{A} is near-zero, where the advantage is defined by the probability of the oracle mistakenly verifying a message:

$$\text{Adv}^{\text{uf-cma}}(\mathcal{A}) = \Pr[\mathcal{V}_F(k, m, t) = 1 \text{ and } m \text{ was not queried to the oracle}]$$

The latter part of the probability lets us ignore **replay attacks** and trivial breaks of the scheme.

A Toy Example

Suppose we take a simple MAC scheme that prepends each message block with a counter, runs this concatenation through a block cipher, and XORs all of the ciphertexts (see [Figure 4.1](#)).

This can be broken easily if we realize that XORs can cancel each other out. Consider

¹ This lone restriction on the adversary is exactly like the one for **IND-CCA**, where it's trivial to get a perfect advantage if you're allowed to decrypt messages you've encrypted.




Figure 4.1: A simple MAC algorithm.

tags for three pairs of messages and what they expand to

$$\begin{aligned}
 T_1 &= \mathcal{MAC}(X_1 \parallel Y_1) &\longrightarrow& \mathcal{E}_K(1 \parallel X_1) \oplus \mathcal{E}_K(2 \parallel Y_1) \\
 T_2 &= \mathcal{MAC}(X_1 \parallel Y_2) &\longrightarrow& \mathcal{E}_K(1 \parallel X_1) \oplus \mathcal{E}_K(2 \parallel Y_2) \\
 T_3 &= \mathcal{MAC}(X_2 \parallel Y_1) &\longrightarrow& \mathcal{E}_K(1 \parallel X_2) \oplus \mathcal{E}_K(2 \parallel Y_1)
 \end{aligned}$$

If we combine these three tags, we can actually derive the tag for a new pair of messages!

$$\begin{aligned}
 T_1 \oplus T_2 \oplus T_3 &= \boxed{\mathcal{E}_K(1 \parallel X_1)} \oplus \boxed{\mathcal{E}_K(2 \parallel Y_1)} \oplus \\
 &\quad \boxed{\mathcal{E}_K(1 \parallel X_1)} \oplus \mathcal{E}_K(2 \parallel Y_2) \oplus \\
 &\quad \mathcal{E}_K(1 \parallel X_2) \oplus \boxed{\mathcal{E}_K(2 \parallel Y_1)} \\
 &= \mathcal{E}_K(2 \parallel Y_2) \oplus \mathcal{E}_K(1 \parallel X_2) && \text{cancel duplicate XORs (highlighted)} \\
 &= \mathcal{MAC}(X_2 \parallel Y_2)
 \end{aligned}$$

Since we haven't queried the tagging algorithm with this particular message, it becomes a valid pairing that breaks the scheme. It's also trivially a reasonable attack, requiring only $q_t = 3$ queries to the tagging algorithm, $\mu = 3$ messages, and the time it takes to perform 3 XORs (if we don't count the internals of \mathcal{MAC}).

4.3 Mode of Operation: CBC-MAC

We state an important fact without proof; it acts as our inspiration for this section:

Theorem 4.1. *Any PRF function yields a UF-CMA secure MAC.*

This means that any secure blockcipher (like [AES](#)) can be used as a MAC. However, they only operate on short input messages. Can we extend our [Modes of Operation](#) to allow MACs on arbitrary-length messages?

Enter CBC-MAC, which looks remarkably like CBC mode for encryption (see [subsection 3.1.2](#)) but disregards all but the last output ciphertext. Given an n -bit block

cipher, $\mathcal{E} : \{0, 1\}^k \times \{0, 1\}^n \mapsto \{0, 1\}^n$, the output message space is $\mathcal{MsgSp} = \{0, 1\}^{mn}$, **fixed** m -block messages (obviously $m \geq 1$).

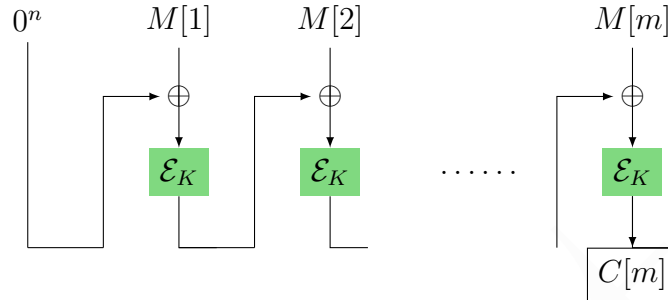


Figure 4.2: The CBC-MAC authentication mode.

To reiterate, this scheme is secure under **UF-CMA** only for a fixed message length across all messages. That is, we can't send messages that are longer or shorter than some predefined multiple of n bits.

Theorem 4.2. *The CBC-MAC authentication scheme is secure if the underlying blockcipher is secure.*

More specifically, for any efficient adversary \mathcal{A} , there exists an adversary \mathcal{B} with similar resources such that \mathcal{A} 's advantage is worse than \mathcal{B} 's:

$$\text{Adv}_{\text{CBC-MAC}}^{\text{uf-cma}}(\mathcal{A}) \leq \text{Adv}_E^{\text{prf}}(\mathcal{B}) + \frac{m^2 q_{\mathcal{A}}^2}{2^{n-1}}$$

(the last term is an artifact of the birthday paradox)

This is an important limitation, and it will be enlightening for the reader to determine why variable-length messages break the CBC-MAC authentication scheme. There *are*, however, ways to extend CBC-MAC to allow variable-length messages, such as by prepending the length as the first message block.

HASH FUNCTIONS

Realize you won't master data structures until you are working on a real-world problem and discover that a hash is the solution to your performance woes.

— Robert Love

An essential part of modern cryptography, **hash functions** transform arbitrary-length input data to a short, fixed-size **digest**:

$$\mathcal{H} : \{0, 1\}^{<2^{64}} \mapsto \{0, 1\}^n$$

Some examples of modern hash functions include those in [Table 5.1](#). They should be pretty familiar: SHA-1 is used by `git` and SHA-3 is used by various cryptocurrencies like Ethereum.¹ They are used as building blocks for encryption, hash-maps, **blockchains**, key-derivation functions, password-storage mechanisms, and more.

Function	Digest Size	Secure?
MD4	128	×
MD5	128	×
SHA-1	160	×
SHA-256	256	✓
SHA-3	224, 256, 384, 512	✓

Table 5.1: A list of some modern hash functions and their output digest length.

¹ Technically, Ethereum uses the KECCAK-256 hash function, which is the pre-standardized version of SHA-3. There are some interesting theories on the difference between the two: though the standardized version changes a [padding rule](#)—allegedly to allow better variability in digest lengths—its underlying algorithm was [weakened to improved performance](#), casting doubts on its general-purpose security.

5.1 Collision Resistance

Not all hash functions are created equal. For example, here's a valid hash function: just output the first n bits of the input as the digest. A **good** hash function tries to distribute its potential inputs uniformly across the output space to minimize *hash collisions*. In fact, most of the functions in Table 5.1 above are considered **broken** from a cryptography perspective: collisions have been found.

Formally, a collision is a pair of messages from the domain, $m_1 \neq m_2$, such that $H(m_1) = H(m_2)$. Obviously, if the domain is larger than the range, there *must* be collisions (by the [pigeonhole principle](#)), but from the perspective of security, we want the probability of *creating* a collision to be very small. This is **collision resistance**.

As we've done several times before, let's formalize the notion of collision resistance with an experiment. If we try to approach this in the traditional way—define an oracle that outputs hashes, and defined some “collision resistance advantage” as the probability of finding two inputs that output the same hash—we immediately run into problems:

1. Since hash algorithms are public, there isn't really a key to keep secret and thus no oracle to construct.
2. Hash functions have collisions *by definition*, so the probability of finding one is *always* one. Even if we, as humans, don't *know* how to find the collision, this is a separate issue.

To get around this, we'll instead consider experiments on *families* of hash functions, where a “key” acts as a selector of specific instances from the family.

This is unfortunate in some ways, because it distances us from concrete hash functions like SHA1. But no alternative is known.

— Introduction to Modern Cryptography (pp. 141)

Formally, we define a family of hash functions as being:

$$\mathcal{H} : \{0, 1\}^k \times \{0, 1\}^m \mapsto \{0, 1\}^n$$

Then, the key is chosen randomly ($k \xleftarrow{\$} \{0, 1\}^k$) and provided to the adversary (to enable actually running the hash functions), who tries to find two inputs that map to the same output.

DEFINITION 5.1: Collision Resistance

A family of hash functions \mathcal{H} is considered **collision resistant** if an adversary's **cr-advantage**—the probability of finding a collision—on a randomly-

chosen instance \mathcal{H}_k is small (≈ 0).

$$\text{Adv}_{\mathcal{H}}^{\text{cg}}(\mathcal{A}) = \Pr[\mathcal{H}_k(x_1) = \mathcal{H}_k(x_2)] \quad \text{where } x_1 \neq x_2$$

This avoids the aforementioned problem of adversaries hard-coding *a priori*-known collisions to specific instances. There's still a bit of a gap between this theoretical security definition and practice, since hash functions still typically don't have keys.

Practice: Find a Collision Do blockciphers make good hash functions? By their very nature (being a **permutation**), their output is collision resistant. However, they don't accept arbitrary-length inputs, and a **mode of operation** will still render arbitrary-length *outputs* while we need a fixed-size digest as a result.

Consider a simple way of combining AES inputs: XOR the individual output blocks. For simplicity, we'll limit ourselves to two AES blocks, so our function family is:

$$\mathcal{H} : \{0, 1\}^k \times \{0, 1\}^{256} \mapsto \{0, 1\}^{128}$$

Is \mathcal{H} collision resistant?

Obviously not. It's actually quite trivial to get the exact same digest, since $x_1 \oplus x_1 = 0$. That is, we pass the same 128-bit block in twice:

$$\begin{aligned} \text{Let } x &\xleftarrow{\$} \{0, 1\}^{128} \text{ and } m = x \parallel x : \\ \mathcal{H}_k(m) &= \text{AES}(x) \oplus \text{AES}(x) \\ &= c \oplus c = 0 \end{aligned}$$

Notice that this is extremely general-purpose, finding 2^{128} messages that all collide to the same value of zero.

5.2 Building Hash Functions

Suppose we had a hash function that compressed short inputs into even-shorter outputs:

$$\mathcal{H}_s : \{0, 1\}^k \times \{0, 1\}^{b+n} \mapsto \{0, 1\}^n$$

where b is relatively small. We can use a technique called the **Merkle-Damgård transform** to create a new compression function that operates on *much* larger inputs, on an arbitrary domain D :

$$\mathcal{H}_\ell : \{0, 1\}^k \times D \mapsto \{0, 1\}^n$$

The algorithm is straightforward and is formalized in [algorithm 5.1](#). It's used by many modern hash function families, including the MD and SHA families. Visually,




Figure 5.1: A visualization of the Merkle-Damgård transform.

ALGORITHM 5.1: The Merkle-Damgård transform, building an arbitrary-length compression function using a limited compression function.

Input: $h(\cdot)$, a limited-range compression function operating on b -bit inputs.

Input: M , the arbitrary-length input message to compress.

Result: M compressed to an n -bit digest.

```

 $m := \|M\|_b$                                 // the number of  $b$ -bit blocks in  $M$ 
 $M[m+1] := \langle M \rangle$                     // the last block is message size
 $V[0] := 0^n$ 
for  $i = 1, \dots, m+1$  do
  |  $V[i] := h(M[i] \| V[i-1])$ 
end
return  $V[m+1]$ 

```

it looks like [Figure 5.1](#): each “block” of the input message is concatenated with the hashed version of its previous block, then hashed again.

The [good news](#) of this transform is that if \mathcal{H}_s is [collision resistant](#), then the larger \mathcal{H}_ℓ is collision resistant as well! That means we can build up complex hash functions from simple primitives, as long as those primitives can make promises about collision resistance. Can they?

Birthday Attacks Recall the [birthday paradox](#): as the number of samples from a range increases, the probability of any two of those samples being equal grows rapidly (there’s a 95% chance that two people at a 50-person party will have the same birthday).

A hash function is **regular** if every range point has the same number of pre-images (that is, if every output has the same number of possible inputs). For such a function, the “birthday attack” finds a collision in $\approx 2^{n/2}$ trials. For a hash function that is *not* regular, such an attack could succeed even sooner.

Thorough research into the modern hash functions (for which $n \geq 160$, large-enough

to protect against birthday attacks) suggests that they are “close to regular.”² Thus, we can safely use them as building blocks for Merkle-Damgård.

Attacks in Practice: SHAttered A collision for the SHA-1 hash was found in February of 2017, breaking the hash function in practice after it was broken theoretically in 2005: two PDFs resolved to the same digest. The attack took $2^{63} - 1$ computations; this is 100,000 faster than the birthday attack.

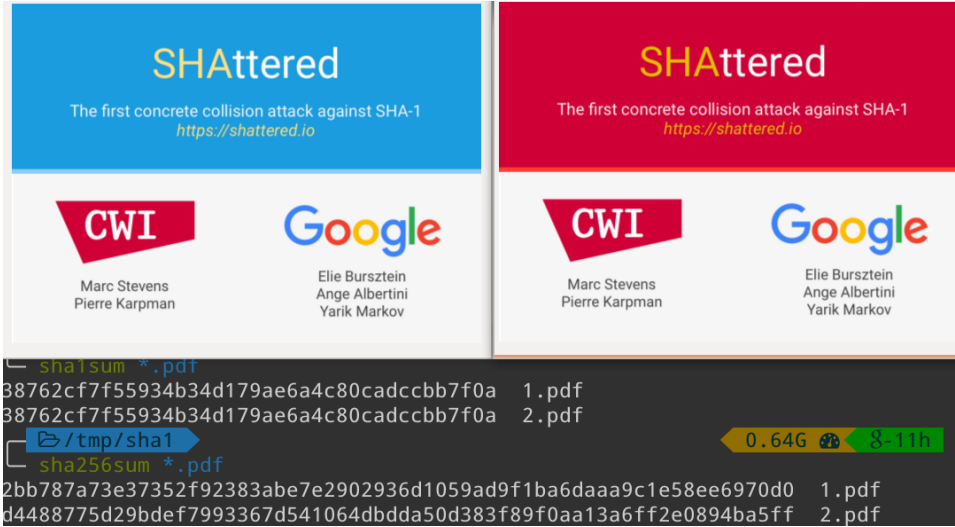


Figure 5.2: The two PDFs in the SHAttered attack and their resulting, identical digests. More details are available on the attack’s [site](https://shattered.io) (because no security attack is complete without a trendy title and domain name).

QUICK MAFFS: Function Nomenclature

Because mathematicians like to use opaque terminology, it’s worth expanding upon the nomenclature for clarity.

The **domain** and **range** of a function should be familiar to us: the domain is a set of inputs and the range (also confusingly called the **codomain** sometimes) is the set of possible outputs for that input. Formally,

$$R = \{f(d) : d \in D\}$$

Example If the domain of $f(x) = x^2$ is all real numbers ($D = \mathbb{R}$), its range is all positive reals $R = \mathbb{R}^+$ (we’ll treat 0 as a positive number for

² Much like the conjecture that AES is PRF secure, this is thus far unproven. As we’ll see later, neither are the security assumptions behind asymmetric cryptography (e.g. “factoring is hard”). Overall, these conjectures on top of conjectures unfortunately do not inspire much confidence in the overall state of security, yet it’s the best we can do.

brevity).

These terms refer to the function as a whole; we chose the input domain, and the range is the corresponding set of outputs. However, it's also useful to examine subsets of the range and ask the inverse question. That is, what's the domain that corresponds to a particular set of values?

Example Given $f(x) = x$ (the diagonal line passing through the origin), for what subset of the domain is $f(x) > 0$? Obviously, when $x > 0$.

This subset is called the **preimage**. Namely, given a subset of the range, $S \subseteq R$, its preimage is the set of inputs that corresponds to it:

$$P = \{x \mid f(x) \in S\}$$

In summary, a preimage of some outputs of a function is **the set of inputs** that **result** in that output.

5.3 One-Way Functions

Hash functions that are viable for cryptographic use must be being **one-way functions**: they must be easy to compute in one direction, but (very) hard to compute in reverse. That is, given a hash, it should be hard to figure out what input resulted in that hash. Informally, a hash function is one-way if, given y and k , it is infeasible to find y 's preimage under h_k .

As usual, we'll define this notion formally with an experiment. Given a hash function family, $\mathcal{H} : \{0, 1\}^k \times D \mapsto \{0, 1\}^n$, we'll have an oracle randomly-select a key and an input value, providing the adversary with its resulting hash and the key (so they can run hash computations).

$$\text{Let } k \xleftarrow{\$} \{0, 1\}^k \text{ and } x \xleftarrow{\$} D : \\ y = \mathcal{H}_k(x)$$

Now, the adversary wins if they can produce a x' for which $\mathcal{H}_k(x') = y$. Note that x' does not need to be x , only in the preimage of y , so this security definition somewhat-includes **collision resistance**.

DEFINITION 5.2: One-Way Function

A family of hash functions \mathcal{H} is considered **one-way** if an adversary's **ow-advantage**—the probability of finding the randomly-chosen input, x , from

its digest y —on a randomly-chosen instance \mathcal{H}_k is small (≈ 0).

$$\text{Adv}_{\mathcal{H}}^{\text{ow}}(\mathcal{A}) = \Pr[\mathcal{H}_k(x') = y]$$

Given our two security properties for a hash function, do either of them imply the other? That is, are either of these true?

$$\begin{aligned} \text{collision resistance} &\implies \text{one-wayness} \\ \text{one-wayness} &\implies \text{collision resistance} \end{aligned}$$

CR \implies OW For functions in general (not necessarily hash functions), collision resistance **does not** imply one-wayness. Consider the trivial identity function: since it’s **one-to-one**, it’s collision resistant by definition, but it’s obviously not one-way. However, for functions that compress their input (e.g. hash functions), it **does** imply it!

OW \implies CR This implication **does not** hold in all cases. We can easily do a “disproof by counterexample”: suppose we have a one-way hash function g . We construct h to hash an n -bit string by delegating to g , sans the last input bit:

$$h(x_1x_2 \cdots x_n) = g(x_1x_2 \cdots x_{n-1})$$

Since g was one-way, h is also one-way. However, it’s obviously not collision resistant, since we know that when given any n -bit input m

$$h(m_1m_2 \cdots m_{n-1}0) = h(m_1m_2 \cdots m_{n-1}1)$$

5.4 Hash-Based MACs

We’ve come full-circle. Can we use a **cryptographically-secure hash function** (a hash function that is both **collision resistant** and **one-way**) to do authentication? Obviously, we can’t use hash functions directly since there is no key.

However, can we somehow include a key within our hash input? More importantly, can we devise a *provably*-secure hash-based **message authentication code**? Enter the aptly-named **HMAC**, visualized below in [Figure 5.3](#).

First, some definitions. H is our hash function instance that maps from an arbitrary domain to an n -bit digest:

$$H : \mathcal{D} \mapsto \{0, 1\}^n$$

Then, we have a secret key K , and we denote $B \geq n/8$ as the *byte*-length of the message block³ The HMAC is computed using a nested structure, mixing the key

³ Typically, $B = 64$ for modern hash functions like MD5, SHA-1, SHA-256, and SHA-512.

with some constants. Namely, we define

$$K_o = \text{opad} \oplus K \parallel 0^{8B-n} \quad K_i = \text{ipad} \oplus K \parallel 0^{8B-n}$$

where opad and ipad are hex-encoded constants:

$$\text{opad} = 0x \underbrace{5C5C5C\dots}_{\text{repeated } B \text{ times}} \quad \text{ipad} = 0x \underbrace{363636\dots}_{\text{repeated } B \text{ times}}$$

Then, the final tag is a simple combination of the transformed keys:

$$\text{Hmac}_K(K) = H(K_o \parallel H(K_i \parallel M))$$

The specific constants are chosen to simplify the proof of security, having no bearing on the security itself.




Figure 5.3: A visualization of the two-tiered structure of HMAC, the standard keyed message authentication code scheme.

HMAC is easy to implement and fast to compute; it is a core part of many standardized cryptographic constructs. Its useful both as a [message authentication code](#) and as a [key-derivation function](#) (which we'll discuss later in asymmetric cryptography).

Theorem 5.1. *HMAC is a PRF assuming that the underlying compression function H is a PRF.*

INDEX OF TERMS

A

AES 15, 24, 38
authenticity 5, 34, 35

B

birthday paradox 27, 30, 31, 39, 43
block cipher ... 15, 24, 26, 27, 30, 34, 36
blockchain 40

C

collision resistance 41, 45
collision resistant 41, 43, 46
confidentiality 5, 35
cr-advantage 41

D

DES 15, 24

E

exclusive-or 11
exhaustive key-search 24

H

hash function 40
HMAC 46

I

impossibility result 13
IND-CCA 31, 34, 36, 37
IND-CCA advantage 32
IND-CPA 20, 21, 25, 26, 31, 34
IND-CPA advantage 20, 22, 28
IND-CPA-cg 22, 28
IND-CPA-cg advantage 23
initialization vector 16, 21, 32
integrity 5, 34, 35

K

Kerckhoff's principle 14, 35
key distribution 7
key-derivation function 47

M

Merkle-Damgård transform 42
message authentication code . 35, 46, 47
mode of operation 16, 22, 24, 36, 42

O

one-time pad 11, 18, 29, 35
one-way 45, 46
ow-advantage 45

P

perfect security 11, 35
PRF advantage 26, 28
PRF secure 26, 30, 31, 34, 44
private key 7
pseudorandom function 25, 30
public key 7

R

random function 34
replay attack 36, 37

S

Shannon-secure 11, 14, 19, 29, 35
symmetric key 6

U

UF-CMA 37, 39
UF-CMA advantage 37